

# TCP Programming

RES, Lecture 2

---

Olivier Liechti

heig-vd

Haute Ecole d'Ingénierie et de Gestion  
du Canton de Vaud

Est-ce que vous seriez capable d'expliquer les différences principales entre le protocole **TCP** et le protocole **UDP**?

Réponse	Moyenne	Total
Oui, bien sûr.	<div><div></div></div> 29%	14
Oui, je pense.	<div><div></div></div> 49%	24
C'est vieux... j'ai des souvenirs très vagues.	<div><div></div></div> 18%	9
Non, pas vraiment.	<div><div></div></div> 4%	2
Total	<div><div></div></div> 100%	49/49

Est-ce que vous avez déjà fait de la programmation réseau en Java?

Réponse	Moyenne	Total
Oui, j'ai déjà implémenté des clients et des serveurs, pour TCP et UDP	<div><div></div></div> 10%	5
Oui, j'ai déjà implémenté une application client-serveur TCP	<div><div></div></div> 33%	16
Non, mais j'ai déjà utilisé la Socket API dans d'autres langages	<div><div></div></div> 33%	16
Non, je ne sais pas du tout comment on implémente des applications client-serveur	<div><div></div></div> 24%	12
Total	<div><div></div></div> 100%	49/49


Est-ce que vous avez déjà utilisé la plate-forme Node.js?

Réponse	Moyenne	Total
Oui, j'ai déjà développé des applications en Node.js et je suis à l'aise.	<div><div></div></div> 6%	3
Oui, j'ai installé Node.js et j'ai fait quelques essais.	<div><div></div></div> 10%	5
Non, j'en ai entendu parler mais je n'ai jamais essayé.	<div><div></div></div> 45%	22
Non, je n'ai jamais entendu parler de Node.js	<div><div></div></div> 39%	19
Total	<div><div></div></div> 100%	49/49

Est-ce que vous connaissez et avez utilisé la technologie **Vagrant**?

Réponse	Moyenne	Total
Oui, j'ai déjà fait des essais avec Vagrant.	<div><div></div></div> 2%	1
Non, mais je sais à quoi sert Vagrant.	<div><div></div></div> 16%	8
Vagrant?	<div><div></div></div> 82%	40
Total	<div><div></div></div> 100%	49/49

# Client-Server Programming





HTTP



SMTP



Spotify®

Proprietary Protocol



# What is an Application-Level Protocol?

---

- **A set of rules** that specify how the application components (e.g. clients and servers) **communicate with each other**. Typically, a protocol defines at least:
  - **Which transport-layer protocol** is used to exchange application-level messages. (e.g. TCP for HTTP)
  - **Which port number(s)** to use (e.g. 80 for HTTP)
  - **What kind of messages** are exchanged by the application components and the **structure** of these messages.
  - The **actions** that need to be taken when these messages are received and the **effect** that is expected.
  - Whether the protocol is **stateful** or **stateless**. In other words, whether the protocol requires the server to manage a session for every connected client.



# Network Programming

---

*Given a application-level protocol,  
how can we implement a client and server in a  
particular programming language?*

***What abstractions, APIs, libraries are  
available to help us do that?***

*We know about TCP, UDP and IP. But how can  
we benefit from these protocols in our code?*

# The TCP Protocol



heig-vd

Haute Ecole d'Ingénierie et de Gestion  
du Canton de Vaud



TCP



UDP

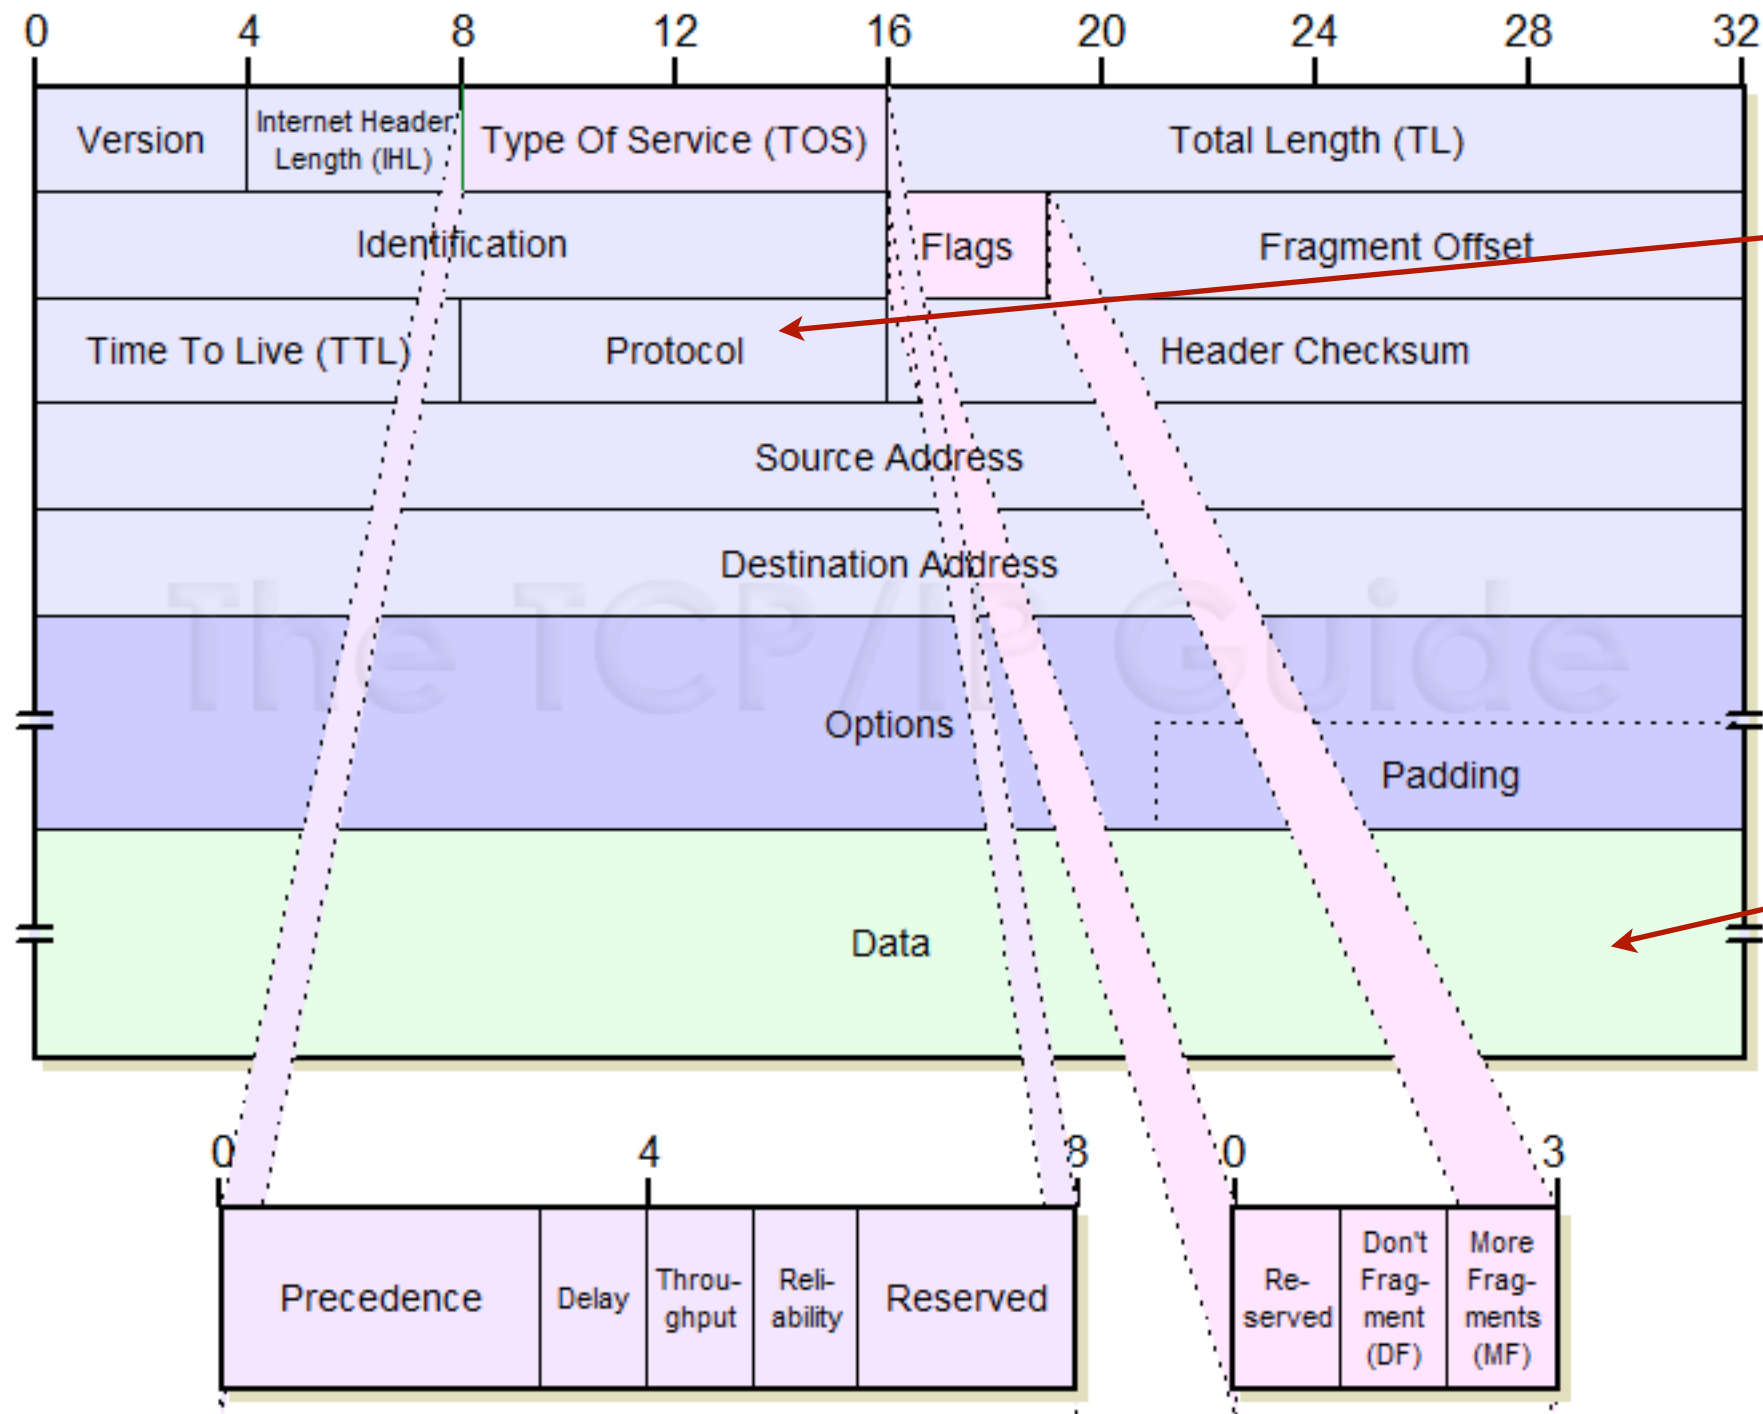


# Transport Protocols

---

- Both TCP and UDP are **transport protocols**.
- This means that they make it possible for **two programs** (i.e. applications, processes) possibly running on **different machines** to **exchange data**.
- The two protocols also make it possible for several programs to **share the same network interface**. They use the notion of **port** for this purpose.
- TCP and UDP define the **structure of messages**. With TCP, messages are called **segments**. With UDP, messages are used **datagrams**.
- The structure of TCP segments (**number and size of headers**) is more complex than the structure of UDP datagrams.
- Both TCP segments and UDP datagrams can be **encapsulated in IP packets**. In that case, we say that the **payload** of the IP packet is a TCP segment, respectively a UDP datagram.

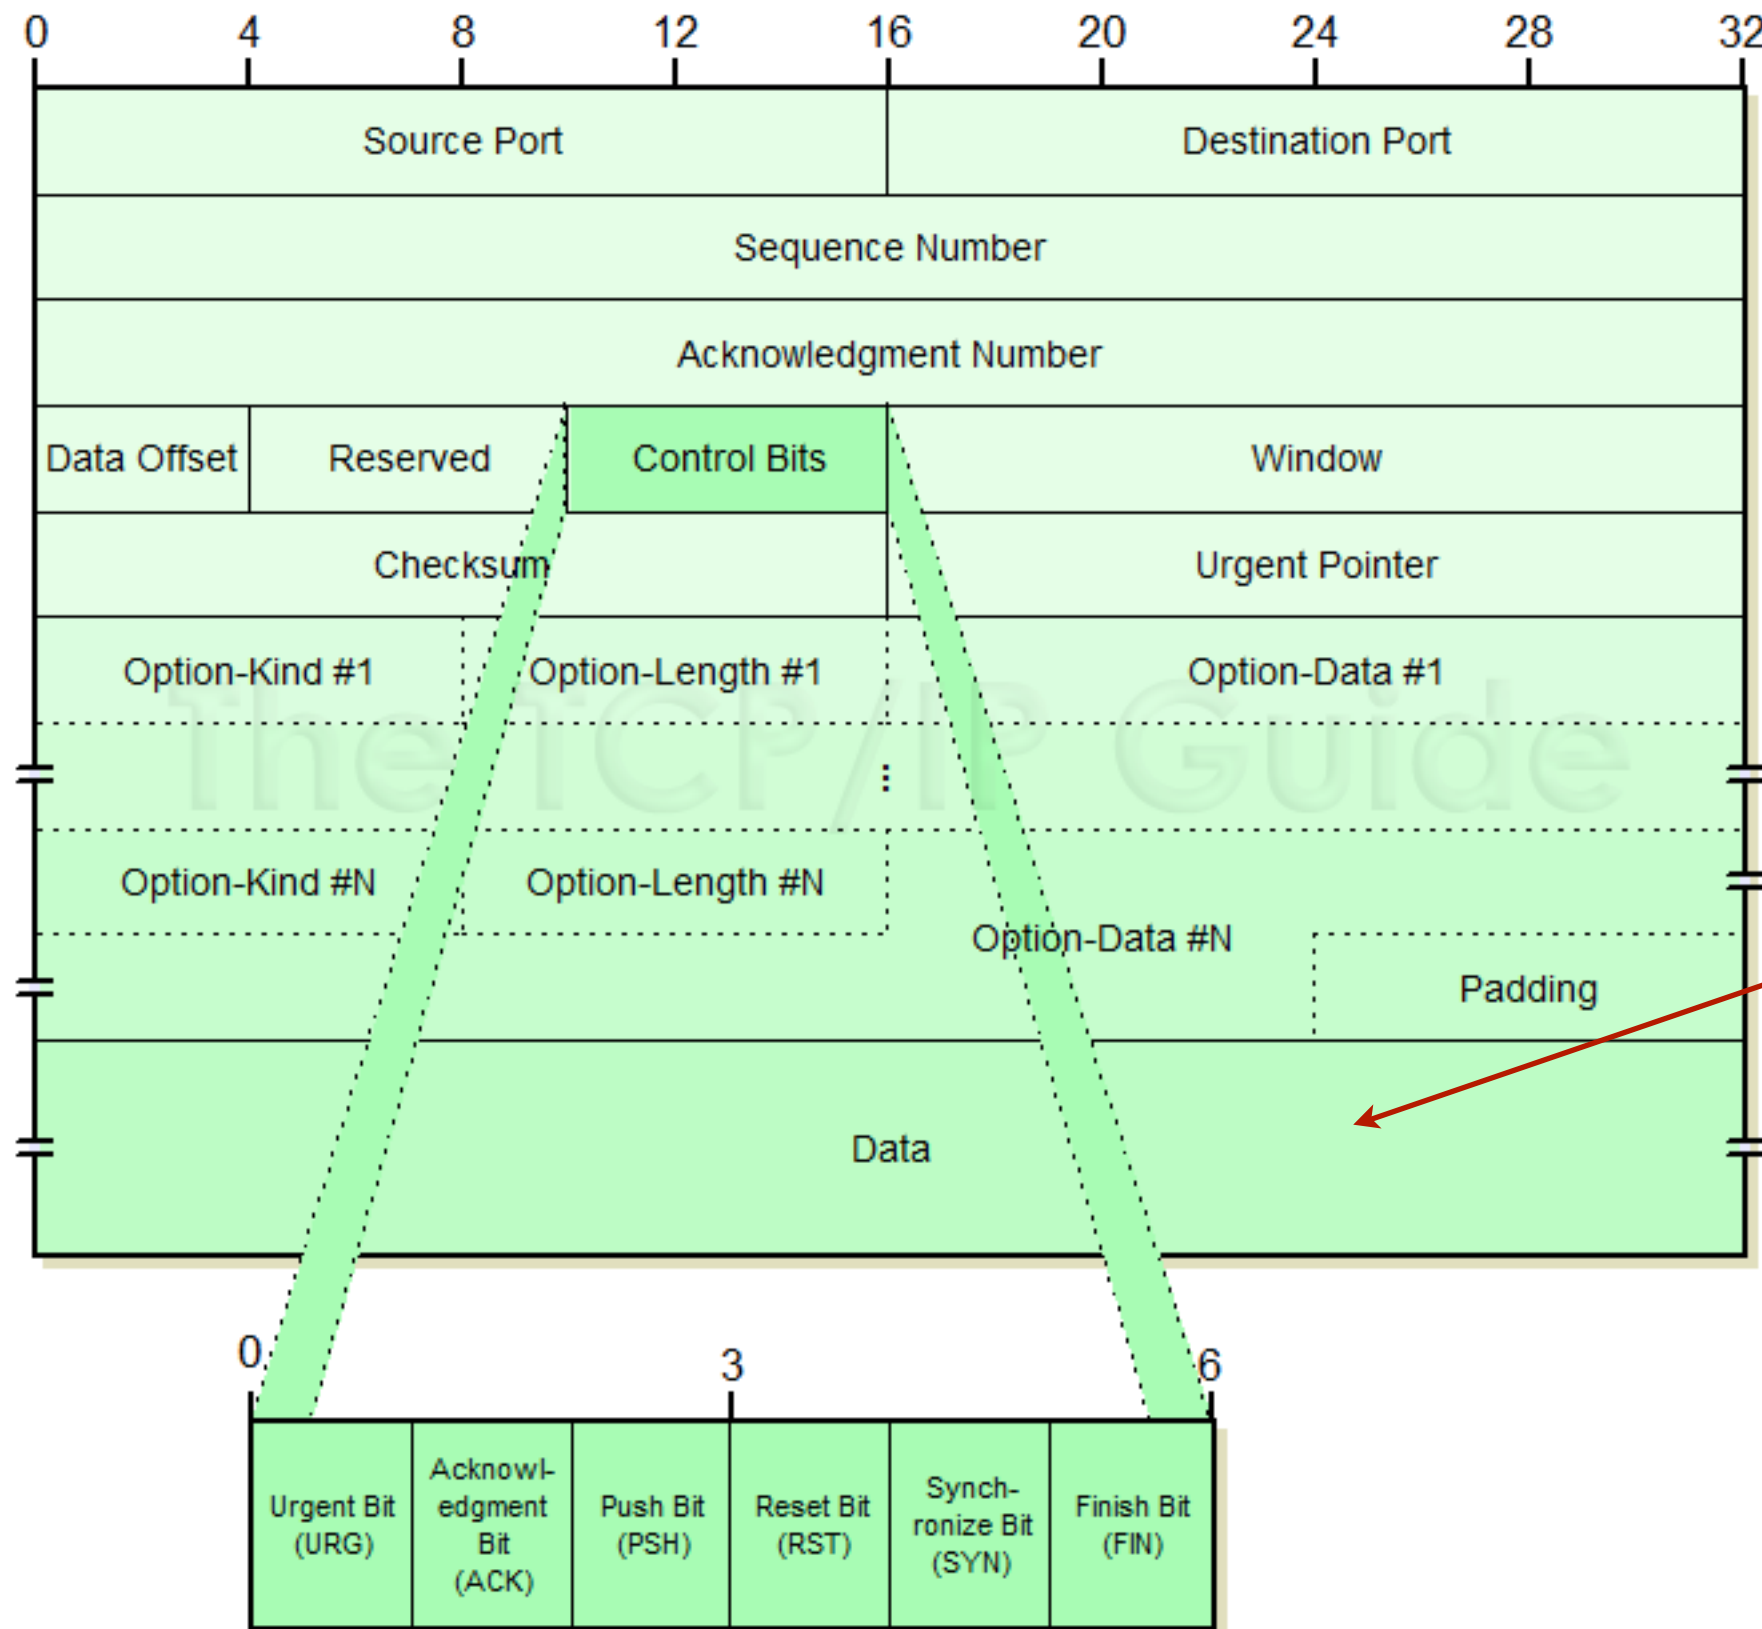
- TCP provides a **connection-oriented service**. The client and the server first have to establish a connection. They can then exchange data through a **bi-directional stream of bytes**.
- TCP provides a **reliable data transfer service**. It makes sure that all bytes sent by one program are received by the other. It also preserves the **ordering** of the exchanged bytes.
- UDP provides a **connectionless service**. The client can send information to the server at any time, **even if there is no server listening**. In that case, the information will simply be lost.
- UDP **does not guarantee the delivery** of datagrams. It is possible that a datagram sent by one client will never reach its destination. The ordering is not guaranteed either.
- TCP supports **unicast** communication. UDP supports **unicast, broadcast and multicast** communication. This is useful for **service discovery**.



If "Data" is a TCP segment, this field has the decimal value "6". If it is a UDP datagram, this field has the decimal value "17".

This can contain a TCP segment, a UDP datagram, or something else.





The bytes that you write in your java program will be here...

Example: **telnet www.heig-vd.ch 80**





# The Socket API



# Network Programming

---

*Given a application-level protocol,  
how can we implement a client and server in a  
particular programming language?*

***What abstractions, APIs, libraries are  
available to help us do that?***

*We know about TCP, UDP and IP. But how can  
we benefit from these protocols in our code?*

# The Socket API


---

- The Socket API is a **standard interface**, which defines **data structures** and **functions** for writing client-server applications.
- It has originally been developed in the context of the Unix operating system and specified as a C API.
- It is now available **across nearly all operating systems and programming environments**.

`<sys/socket.h>`

```

int  accept(int socket, struct sockaddr *address,
          socklen_t *address_len);
int  bind(int socket, const struct sockaddr *address,
          socklen_t address_len);
int  connect(int socket, const struct sockaddr *address,
            socklen_t address_len);
int  getpeername(int socket, struct sockaddr *address,
                socklen_t *address_len);
int  getsockname(int socket, struct sockaddr *address,
                 socklen_t *address_len);
int  getsockopt(int socket, int level, int option_name,
                void *option_value, socklen_t *option_len);
int  listen(int socket, int backlog);
ssize_t recv(int socket, void *buffer, size_t length, int flags);
ssize_t recvfrom(int socket, void *buffer, size_t length,
                  int flags, struct sockaddr *address, socklen_t *address_len);
ssize_t recvmsg(int socket, struct msghdr *message, int flags);
ssize_t send(int socket, const void *message, size_t length, int flags);
ssize_t sendmsg(int socket, const struct msghdr *message, int flags);
ssize_t sendto(int socket, const void *message, size_t length, int flags,
               const struct sockaddr *dest_addr, socklen_t dest_len);
int  setsockopt(int socket, int level, int option_name,
                const void *option_value, socklen_t option_len);
int  shutdown(int socket, int how);
int  socket(int domain, int type, int protocol);
int  socketpair(int domain, int type, int protocol,
                int socket_vector[2]);
  
```



# Using the Socket API for a TCP **Server**



1. Create a "receptionist" **socket**
2. **Bind** the socket to an IP address / port
3. Loop
  - 3.1. **Accept** an incoming connection (**block** until a client arrives)
  - 3.2. Receive a new socket when a client has arrived
  - 3.3. **Read** and **write** bytes through this socket, communicating with the client
  - 3.4. **Close** the client socket (and go back to listening)
4. **Close** the "receptionist" socket

# Using the Socket API for a TCP **Client**


---

1. Create a **socket**
2. Make a **connection request** on an IP address / port
3. **Read** and **write** bytes through this socket, communicating with the client
4. **Close** the client socket

# Using the Socket API



# Using the Socket API in Java

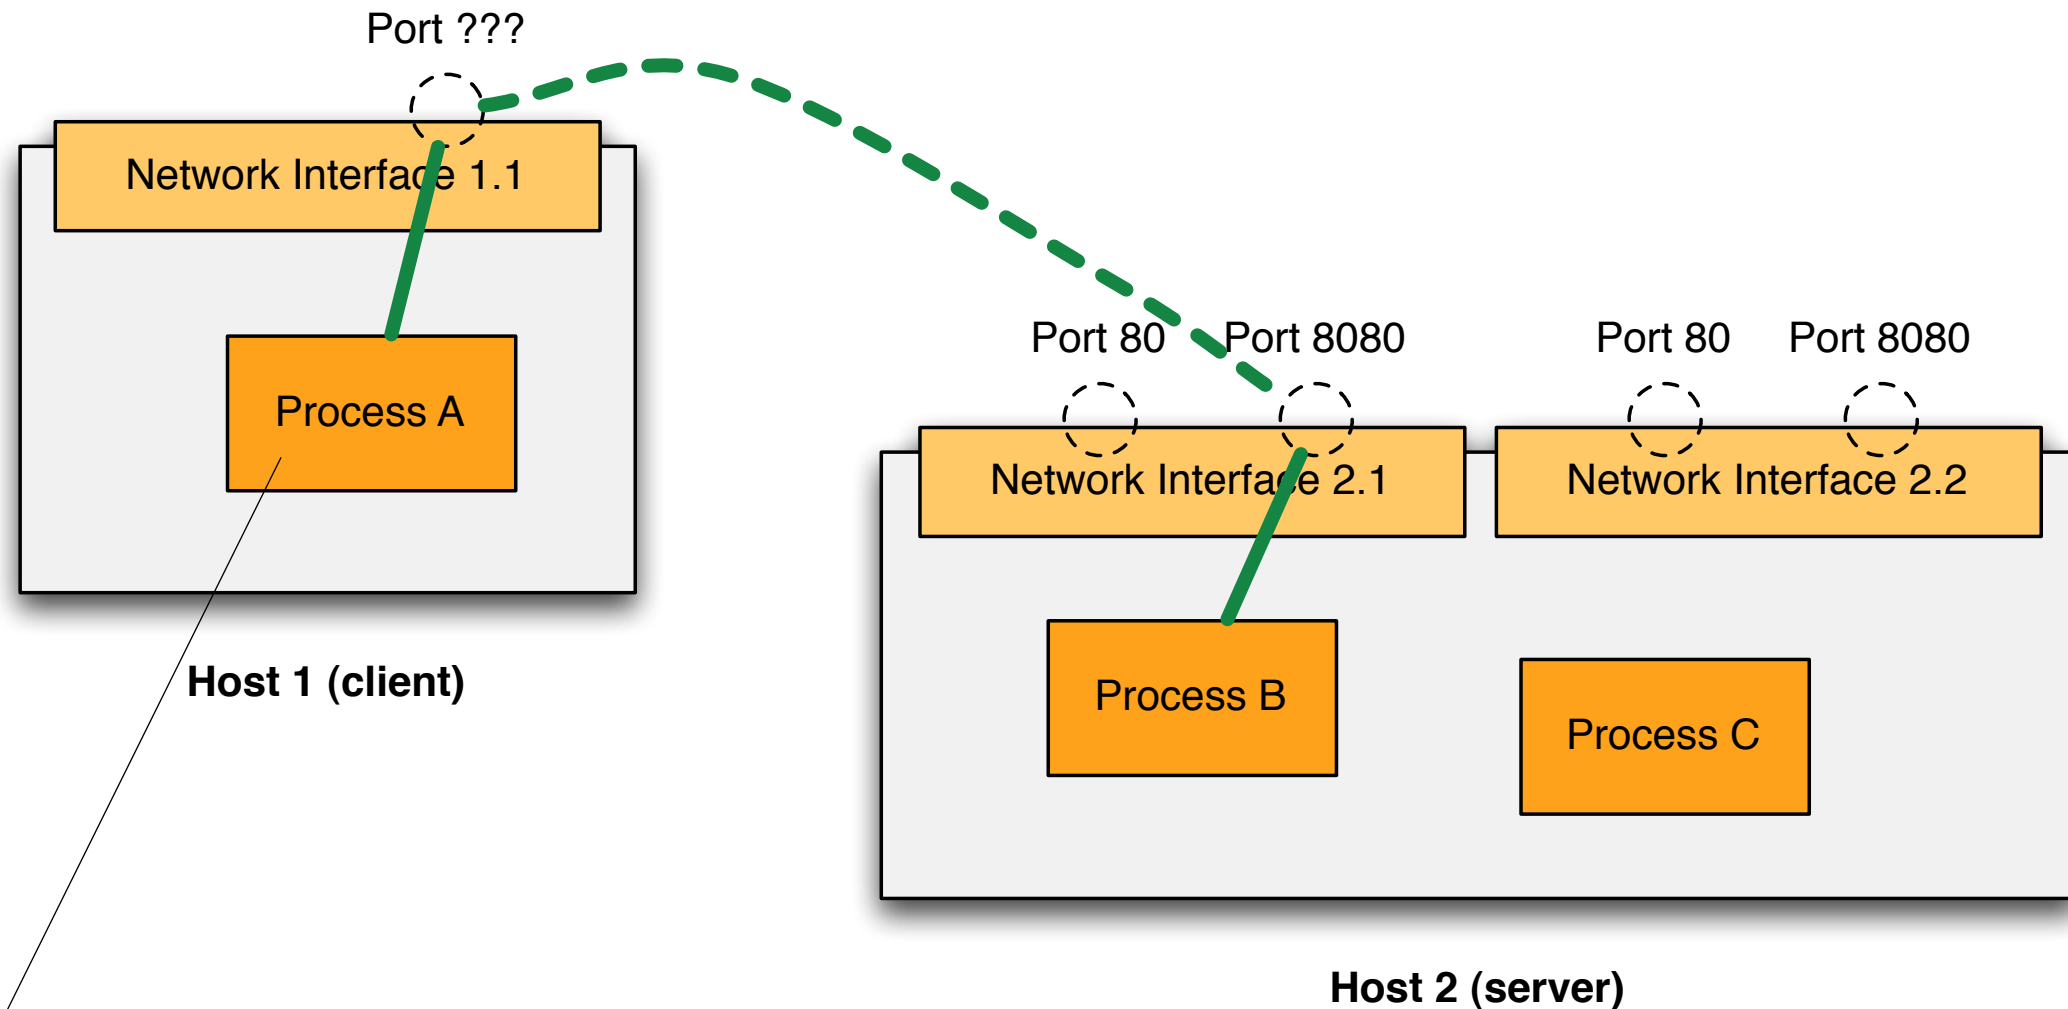


```
// Listen on port 8080
ServerSocket serverSocket = new ServerSocket(8080);

// Wait (block) until a client makes a connection request...
Socket commSocket = serverSocket.accept();
```




# Using the Socket API in Java



```
// Makes a connection request on port 8080
Socket serverSocket = new Socket("host2", 8080);
```

# Using the Socket API in Java



```
// Makes a connection request on port 8080
Socket socket = new Socket("host2", 8080);
```

```
InputStream fromServer = socket.getInputStream();
OutputStream toServer = socket.getOutputStream();
```

```
// Listen on port 8080
ServerSocket serverSocket = new ServerSocket(8080);
```

```
// Wait until a client makes a connection request...
Socket commSocket = serverSocket.accept();
```

```
InputStream fromClient = commSocket.getInputStream();
OutputStream toClient = commSocket.getOutputStream();
```

Example: **05-DumbHttpClient**



# Code walkthrough

```
public void sendWrongHttpRequest() {  
    Socket clientSocket = null;  
    OutputStream os = null;  
    InputStream is = null;
```

```
    try {
```

```
        clientSocket = new Socket("www.heig-vd.ch", 80);
```

```
        os = clientSocket.getOutputStream();
```

```
        is = clientSocket.getInputStream();
```

**get streams to send and  
receive bytes**

```
        String malformedHttpRequest = "Hello, sorry, but I don't speak HTTP...\r\n\r\n";  
        os.write(malformedHttpRequest.getBytes());
```

```
        ByteArrayOutputStream responseBuffer = new ByteArrayOutputStream();
```

```
        byte[] buffer = new byte[BUFFER_SIZE];
```

```
        int newBytes;
```

```
        while ((newBytes = is.read(buffer)) != -1) {  
            responseBuffer.write(buffer, 0, newBytes);  
        }
```

**read bytes sent by the  
server until the  
connection is closed**

```
        LOG.log(Level.INFO, "Response sent by the server: ");
```

```
        LOG.log(Level.INFO, responseBuffer.toString());
```

```
    } catch (IOException ex) {
```

```
        LOG.log(Level.SEVERE, null, ex);
```

```
    } finally {
```

```
    }
```

```
}
```

...

**establish a connection  
with server**

Example: **04-StreamingTimeServer**



# Code walkthrough

```
ServerSocket serverSocket = null;  
Socket clientSocket = null;  
BufferedReader reader = null;  
PrintWriter writer = null;
```

```
try {
```

```
    serverSocket = new ServerSocket(listenPort);
```

```
    logServerSocketAddress(serverSocket);
```

```
    clientSocket = serverSocket.accept();
```

```
    logSocketAddress(clientSocket);
```

```
    reader = new BufferedReader(new InputStreamReader(clientSocket.getInputStream()));
```

```
    writer = new PrintWriter(clientSocket.getOutputStream());
```

```
    for (int i = 0; i < numberOfIterations; i++) {  
        writer.println(String.format("{'time' : '%s'}", new Date()));
```

```
        writer.flush();
```

```
        LOG.log(Level.INFO, "Sent data to client, doing a pause...");
```

```
        Thread.sleep(pauseDuration);
```

```
    }
```

```
    } catch (IOException | InterruptedException ex) {
```

```
        LOG.log(Level.SEVERE, ex.getMessage());
```

```
    } finally {
```

```
        reader.close();
```

```
        writer.close();
```

```
        clientSocket.close();
```

```
        serverSocket.close();
```

```
    }
```

**bind on TCP port**

**block until a client makes a  
connection request**

**we want to exchange  
characters with the  
clients (we should  
specify the encoding!)**

**we make sure to flush  
the buffer, so that  
characters are actually  
sent!**

# Introduction to Vagrant



# How to use Vagrant?

---

- Vagrant was initially created to **manage Virtual Box VMs**. Today, other types of VMs (on local machines and in the cloud) are supported.
- Essentially, the idea is that instead of using the Virtual Box GUI to create, configure, control and use your VMs, you write scripts and use command line tools.
- “**Provisioning**” is the process of installing additional software on top of a “box”. There are different ways to do that: shell scripts and DevOps tools such as Puppet, Chef or Ansible.
- The **community is sharing “boxes”**, which you can use as a starting point.





Discover Vagrant Boxes

https://atlas.hashicorp.com/boxes/search?utf8=✓&sort=&provider=&q=jenkins

ATLAS

Sign up

Sign in

Discover Vagrant Boxes

This page lets you discover and use Vagrant Boxes created by the community. You can search by operating system, architecture or provider.

jenkins

Provider filter

virtualbox

vmware\_desktop

digitalocean

aws

rackspace

hyperv

parallels


Sort by

Recently Created

Recently Updated


Downloads

Favorites




[scottpgallagher/ubuntu-14\\_04-jenkins](#) Ubuntu 14.04 with Jenkins and Docker

411 downloads | 0.1.0 | last release 8 months ago




[wgarcia/centos65-jenkins](#) Jenkins on CentOS 6.5

289 downloads | 0.1.0 | last release 8 months ago




[jnj-devops/centos65-jenkins](#) Centos 6.5 and Jenkins 1.576 devbox.

37 downloads | 0.1.0 | last release 6 months ago



[kristallizer/adobe-jenkins](#)

35 downloads | 0.1.1 | last release 3 months ago



[dzuluaga/apigee-ci-jenkins-maven-git-jmeter](#)

17 downloads | 0.1.0 | last release 11 months ago

Previous

Next

Help

Terms

Privacy

Security

© 2015 HashiCorp, Inc.

# How to use Vagrant?

Usage: vagrant [options] <command> [<args>]

-v, --version	Print the version and exit.
-h, --help	Print this help.

Common commands:

<b>box</b>	<b>manages boxes: installation, removal, etc.</b>
connect	connect to a remotely shared Vagrant environment
destroy	stops and deletes all traces of the vagrant machine
global-status	outputs status Vagrant environments for this user
halt	stops the vagrant machine
help	shows the help for a subcommand
<b>init</b>	<b>initializes a new Vagrant environment by creating a Vagrantfile</b>
login	log in to Vagrant Cloud
package	packages a running vagrant environment into a box
plugin	manages plugins: install, uninstall, update, etc.
<b>provision</b>	<b>provisions the vagrant machine</b>
rdp	connects to machine via RDP
reload	restarts vagrant machine, loads new Vagrantfile configuration
resume	resume a suspended vagrant machine
share	share your Vagrant environment with anyone in the world
<b>ssh</b>	<b>connects to machine via SSH</b>
ssh-config	outputs OpenSSH valid configuration to connect to the machine
status	outputs status of the vagrant machine
suspend	suspends the machine
<b>up</b>	<b>starts and provisions the vagrant environment</b>
version	prints current and latest Vagrant version

# The Vagrantfile

what kind of “base” box  
do we want to use

```
# -*- mode: ruby -*-  
# vi: set ft=ruby :
```

```
# Vagrantfile API/syntax version. Don't touch unless you know what you're doing!  
VAGRANTFILE_API_VERSION = "2"
```

```
Vagrant.configure(VAGRANTFILE_API_VERSION) do |config|
```

```
  config.vm.box = "phusion/ubuntu-14.04-amd64"
```

```
  config.vm.network "private_network", ip: "192.168.42.42"
```

the IP address assigned  
to the VM

```
  config.vm.provision "shell", path: "provision.sh", privileged: false
```

```
  # config.vm.network "forwarded_port", guest: 80, host: 8080
```

```
  # config.ssh.forward_agent = true
```

```
  # config.vm.synced_folder "../data", "/vagrant_data"
```

```
  config.ssh.forward_x11 = true
```

```
  # config.vm.provider "virtualbox" do |vb|
```

```
    # Don't boot with headless mode
```

```
    vb.gui = true
```

```
  #
```

```
  # Use VBoxManage to customize the VM. For example to change memory:
```

```
  vb.customize ["modifyvm", :id, "--memory", "1024"]
```

```
  # end
```

```
end
```

this script is executed at  
“provisioning time”

PAT: localhost:8080 will be  
forwarded to  
192.168.42.42:80

(additional) shared folders; by default, the host  
directory containing the Vagrantfile is available  
under /vagrant in the VM

```
config.vm.network "private_network", ip: "192.168.42.42"  
config.vm.network "forwarded_port", guest: 9907, host: 4207
```

heig-vd

Haute Ecole d'Ingénierie et de Gestion  
du Canton de Vaud

# Network

Another laptop (10.192.95.122)

```
$ telnet 10.192.95.121 4207
```

Your laptop (10.192.95.121)

```
$ telnet 192.168.42.42 9907
```

```
$ telnet localhost 4207
```

Vagrant box (192.168.42.42)

→ **listen(9907)**

```
$ telnet localhost 9907
```

# Demo & setup

---

1. Clone the repo (that should be quick)

```
$ git clone git@github.com:SoftEng-HEIGVD/  
Teaching-HEIGVD-RES-2015-Vagrant.git  
$ cd Teaching-HEIGVD-RES-2015-Vagrant/box
```

2. Ask Vagrant to setup and launch the VM

```
$ vagrant up
```

*... wait for the download of the “base” box... (stored in cache)*  
*... wait for the execution of the provision.sh script*

3. Connect to the VM and enjoy a UNIX environment

```
$ vagrant ssh
```

# Demo & setup

---

4. Have a look at the shared folder (files stored on host)

```
$ cd /vagrant  
$ ls -l
```

5. move to the examples folder; build and run with maven

```
$ mvn clean install
```

# La commande **lsof**

## NAME

**lsof** - list open files

## SYNOPSIS

```
lsof [ -?abChKlnNOPRtUvVX ] [ -A A ] [ -c c ] [ +c c ] [ +|-d d ] [ +|-D D ] [ +|-e s ] [ +|-f [cfgGn] ] [ -F [f] ] [ -g [s] ] [ -i [i] ] [ -k k ] [ +|-L [l] ] [ +|-m m ] [ +|-M ] [ -o [o] ] [ -p s ] [ +|-r [t[m<fmt>]] ] [ -s [p:s] ] [ -S [t] ] [ -T [t] ] [ -u s ] [ +|-w ] [ -x [f1] ] [ -z [z] ] [ -Z [Z] ] [ -- ] [names]
```

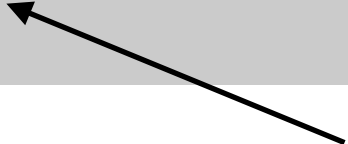
## DESCRIPTION

Lsof revision 4.86 lists on its standard output file information about files opened by processes for the following UNIX dialects:

- Apple Darwin 9 and Mac OS X 10.[567]
- FreeBSD 4.9 and 6.4 for x86-based systems
- FreeBSD 8.2, 9.0 and 10.0 for AMD64-based systems
- Linux 2.1.72 and above for x86-based systems
- Solaris 9, 10 and 11

(See the DISTRIBUTION section of this manual page for information on how to obtain the latest lsof revision.)

An open file may be a regular file, a directory, a block special file, a character special file, an executing text reference, a library, a stream or a network file (Internet socket, NFS file or UNIX domain socket.) A specific file or all the files in a file system may be selected by path.



```
$ lsof
```

```
$ lsof -i | grep -i listen
```

```
$ lsof -i -P | grep -i listen
```



# Handling Concurrency





# Concurrency in Network Programming

---

***You don't want your server to talk to only one client at the time, do you?***

*Even for **stateless** protocols...*



## blocking IO (synchronous)

n employee = n threads

employees are expensive

limited space for employees in the truck

few employees => long queue



## non-blocking IO (asynchronous)

there is only 1 employee (1 thread)

customers are called back when the request is fulfilled

no queue

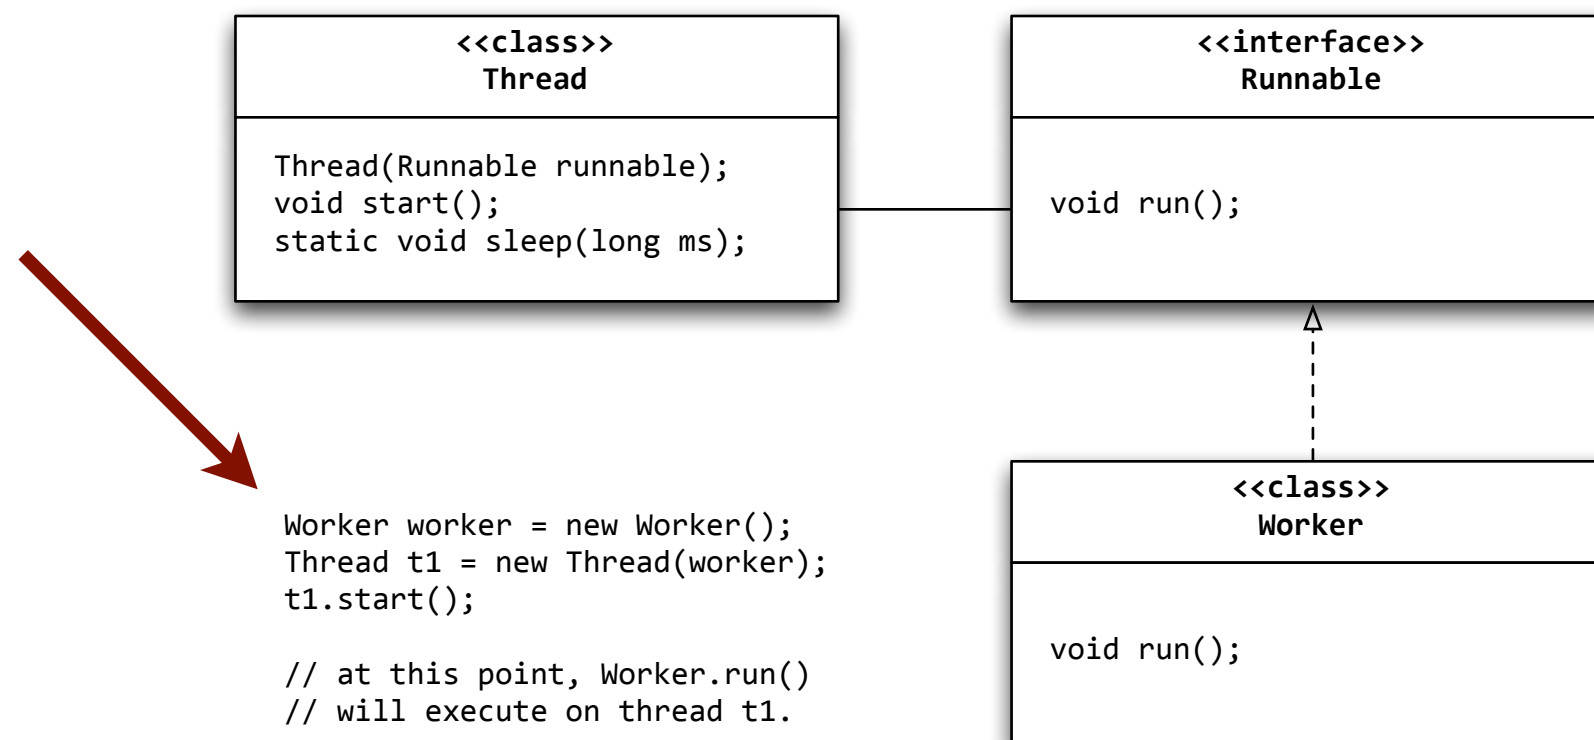
# Concurrent Programming

---

- On top of the **operating system**, it is possible to launch the Java Virtual Machine (**JVM**) several times (by invoking the java command). In this scenario, there is **one process (program) for every JVM instance**.
- If you don't do anything special, there is a **single execution thread** within each JVM. This means that all instructions in your code are executed **sequentially**.
- Very often, you write software where you want to **perform several tasks at the same time** (concurrently). For instance:
  - Manage a UI **while** fetching data from the network,
  - Talking to one HTTP client **while** talking to another HTTP client,
  - Have a worker do complex calculations on a subset of the data, **while** having another worker do the same calculations on another subset.
- You can use **threads** (also called **lightweight processes**) for this purpose.

# Concurrent Programming in Java

- In Java, there are two main types
  - The **Thread class**, which *could be extended* to implement the behavior you want to run in parallel.
  - The **Runnable interface**, which *is implemented* for the same purpose and is passed as an argument to the Thread constructor.



# Concurrent Programming in Java

- There are other classes related to threads, in the `java.util.concurrent` package. An important one is the **ExecutorService**, which makes it possible to use **thread pools**.
- A thread pool gives you a way to limit the number of threads spawned (by your server), so that you will not consume all resources. Others are queued.

[http://www.vogella.com/  
tutorials/  
JavaConcurrency/  
article.html](http://www.vogella.com/tutorials/JavaConcurrency/article.html)

```
package de.vogella.concurrency.threadpools;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
public class Main {
    private static final int NTHREDS = 10;

    public static void main(String[] args) {
        ExecutorService executor = Executors.newFixedThreadPool(NTHREDS);
        for (int i = 0; i < 500; i++) {
            Runnable worker = new MyRunnable(10000000L + i);
            executor.execute(worker);
        }
        // This will make the executor accept no new threads
        // and finish all existing threads in the queue
        executor.shutdown();
        // Wait until all threads are finish
        executor.awaitTermination();
        System.out.println("Finished all threads");
    }
}
```

## Single Threaded Single Process Blocking

**Not really an option...**

The server implements a loop.

It waits for a client to arrive.  
Then services the client until done.

Then only goes back to accept the next  
client.

Can only talk to 1 client at the time

It is only when we reach this line that  
a new client can connect

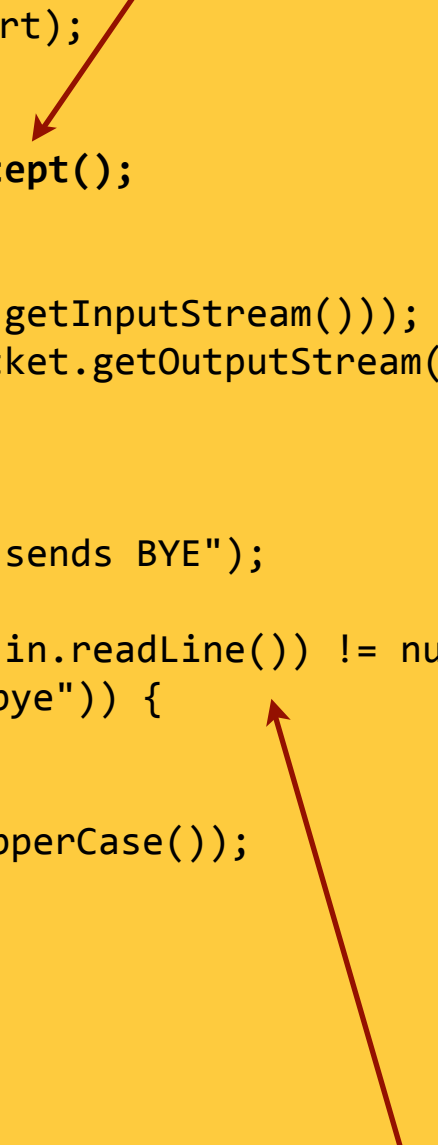
```
serverSocket = new ServerSocket(port);
while (true) {

    clientSocket = serverSocket.accept();

    in = new BufferedReader(new
        InputStreamReader(clientSocket.getInputStream()));
    out = new PrintWriter(clientSocket.getOutputStream());
    String line;
    boolean shouldRun = true;

    LOG.info("Reading until client sends BYE");

    while ( (shouldRun) && (line = in.readLine()) != null ) {
        if (line.equalsIgnoreCase("bye")) {
            shouldRun = false;
        }
        out.println("> " + line.toUpperCase());
        out.flush();
    }
    clientSocket.close();
    in.close();
    out.close();
}
```



It takes a long time to serve each client

## Single Threaded Multi Process Blocking

### How apache httpd did it (with pre-fork, kind of...)

The server implements a loop.  
It waits for a client to arrive.  
When the client arrives, the server forks  
a new process.

The child process serves the client while  
the server is immediately ready to  
serve the next client.

Forking a process is kind of heavy...  
and resource hungry

While the child process serves the client...

... the parent can immediately welcome the next client.

```
while(1) { // main accept() loop
    sin_size = sizeof their_addr;
    new_fd = accept(sockfd, (struct sockaddr *)&their_addr,
&sin_size);
    if (new_fd == -1) {
        perror("accept");
        continue;
    }

    inet_ntop(their_addr.ss_family,
get_in_addr((struct sockaddr *)&their_addr),
s, sizeof s);
    printf("server: got connection from %s\n", s);

    if (!fork()) { // this is the child process
        close(sockfd); // child doesn't need the listener
        if (send(new_fd, "Hello, world!", 13, 0) == -1)
            perror("send");
        close(new_fd);
        exit(0);
    }
    close(new_fd); // parent doesn't need this
}
```



## Multi Threaded Single Process Blocking

### The 'old' Java way

The server uses a first thread to wait for connection requests from clients.

Each time a client arrives, a new thread is created and used to serve the client.

Millions of clients, millions of threads?

Resource hungry.  
Not scalable.

The ReceptionistWorker implements a run() method that will execute on its own thread.

```
private class ReceptionistWorker implements Runnable {

    @Override
    public void run() {
        ServerSocket serverSocket;

        try {
            serverSocket = new ServerSocket(port);
        } catch (IOException ex) {
            LOG.log(Level.SEVERE, null, ex);
            return;
        }

        while (true) {
            LOG.log(Level.INFO, "Waiting for a new client");
            try {
                Socket clientSocket = serverSocket.accept();
                LOG.info("A new client has arrived...");
                new Thread(new ServantWorker(clientSocket)).start();
            } catch (IOException ex) {
                LOG.log(Level.SEVERE, ex.getMessage(), ex);
            }
        }
    }
}
```

As soon as a client is connected, a new thread is created.  
The code that manages the interaction with the client executes on this thread.

**2 types of workers, n+1 threads**



Example: **07-TcpServers**



Single Thread  
Single Process  
Asynchronous Programming

## The 'à la Node.js' way

The server uses a single thread, but in a non-blocking, asynchronous way.

Callback functions have to be written, so that they can be invoked when clients arrive, when data is received, etc.

Different programming logic.  
Scalable.

We are registering callback functions on the various types of events that can be notified by the server...

```
// let's create a TCP server
var server = net.createServer();

// it reacts to events: 'listening', 'connection', 'close', etc.
// register callback functions, to be invoked when the events
// occur (everything happens on the same thread)

server.on('listening', callbackFunctionToCallWhenSocketIsBound);
server.on('connection',
  callbackFunctionToCallWhenNewClientHasArrived);

//Start listening on port 9907
server.listen(9907);

// This callback is called when the socket is bound and is in
// listening mode. We don't need to do anything special.
function callbackFunctionToCallWhenSocketIsBound() {
  console.log("The socket is bound and listening");
  console.log("Socket value: %j", server.address());
}

// This callback is called after a client has connected.
function callbackFunctionToCallWhenNewClientHasArrived(socket) {
  ...
}
```

... and we code these functions, implementing the behavior that is expected when the events occur.

Select is a blocking operation (with a possible timeout). It blocks until something has happened on one of the provided sets of file descriptors.

## Single Thread Single Process IO Multiplexing

### The 'select' way

Sockets are set in a non-blocking state, which means that `read()`, `write()` and other functions do not block.

System calls such as `select()` or `poll()` block, but work on multiple sockets. They return if data has arrived on at least one of the sockets.

Watch out for performance.

```
#include <stdio.h>
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>

int main(void) {
    fd_set rfds;
    struct timeval tv;
    int retval;

    /* Watch stdin (fd 0) to see when it has input. */
    FD_ZERO(&rfds);
    FD_SET(0, &rfds);
    /* Wait up to five seconds. */
    tv.tv_sec = 5;
    tv.tv_usec = 0;

    retval = select(1, &rfds, NULL, NULL, &tv);
    /* Don't rely on the value of tv now! */

    if (retval == -1)
        perror("select()");
    else if (retval)
        printf("Data is available now.\n");
        /* FD_ISSET(0, &rfds) will be true. */
    else
        printf("No data within five seconds.\n");

    return 0;
}
```

Here, we know that something has happened on one of the sockets. We can iterate over the set of file descriptors and get the data.

Example: **QuizRouletteServer**



# Example: **06-PresenceApplication**



# Lab-02: Writing JUnit Tests



- When you fork/clone the lab-02 repo, you will get:
  - **3 maven projects** (-build, -code, -test); you should be familiar with the structure by now, as it is the same as in lab-00 and lab-01.
- The **code project** contains:
  - a working server-side implementation for the Roulette Protocol v1.
  - only the interface for the client-side of the Roulette Protocol v1.
- The **test project** contains
  - Tests to validate the behavior the Roulette Protocol v1 client and server.
  - No tests to validate the implementation of the Roulette Protocol v2.

# The Client Interface

```
public interface IRouletteV1Client {  
  
    public void connect(String server, int port) throws IOException;  
    public void disconnect() throws IOException;  
    public boolean isConnected();  
    public void loadStudent(String fullname) throws IOException;  
    public void loadStudents(List<Student> students) throws IOException;  
    public Student pickRandomStudent() throws EmptyStoreException, IOException;  
    public int getNumberOfStudents() throws IOException;  
    public String getProtocolVersion() throws IOException;  
  
}
```

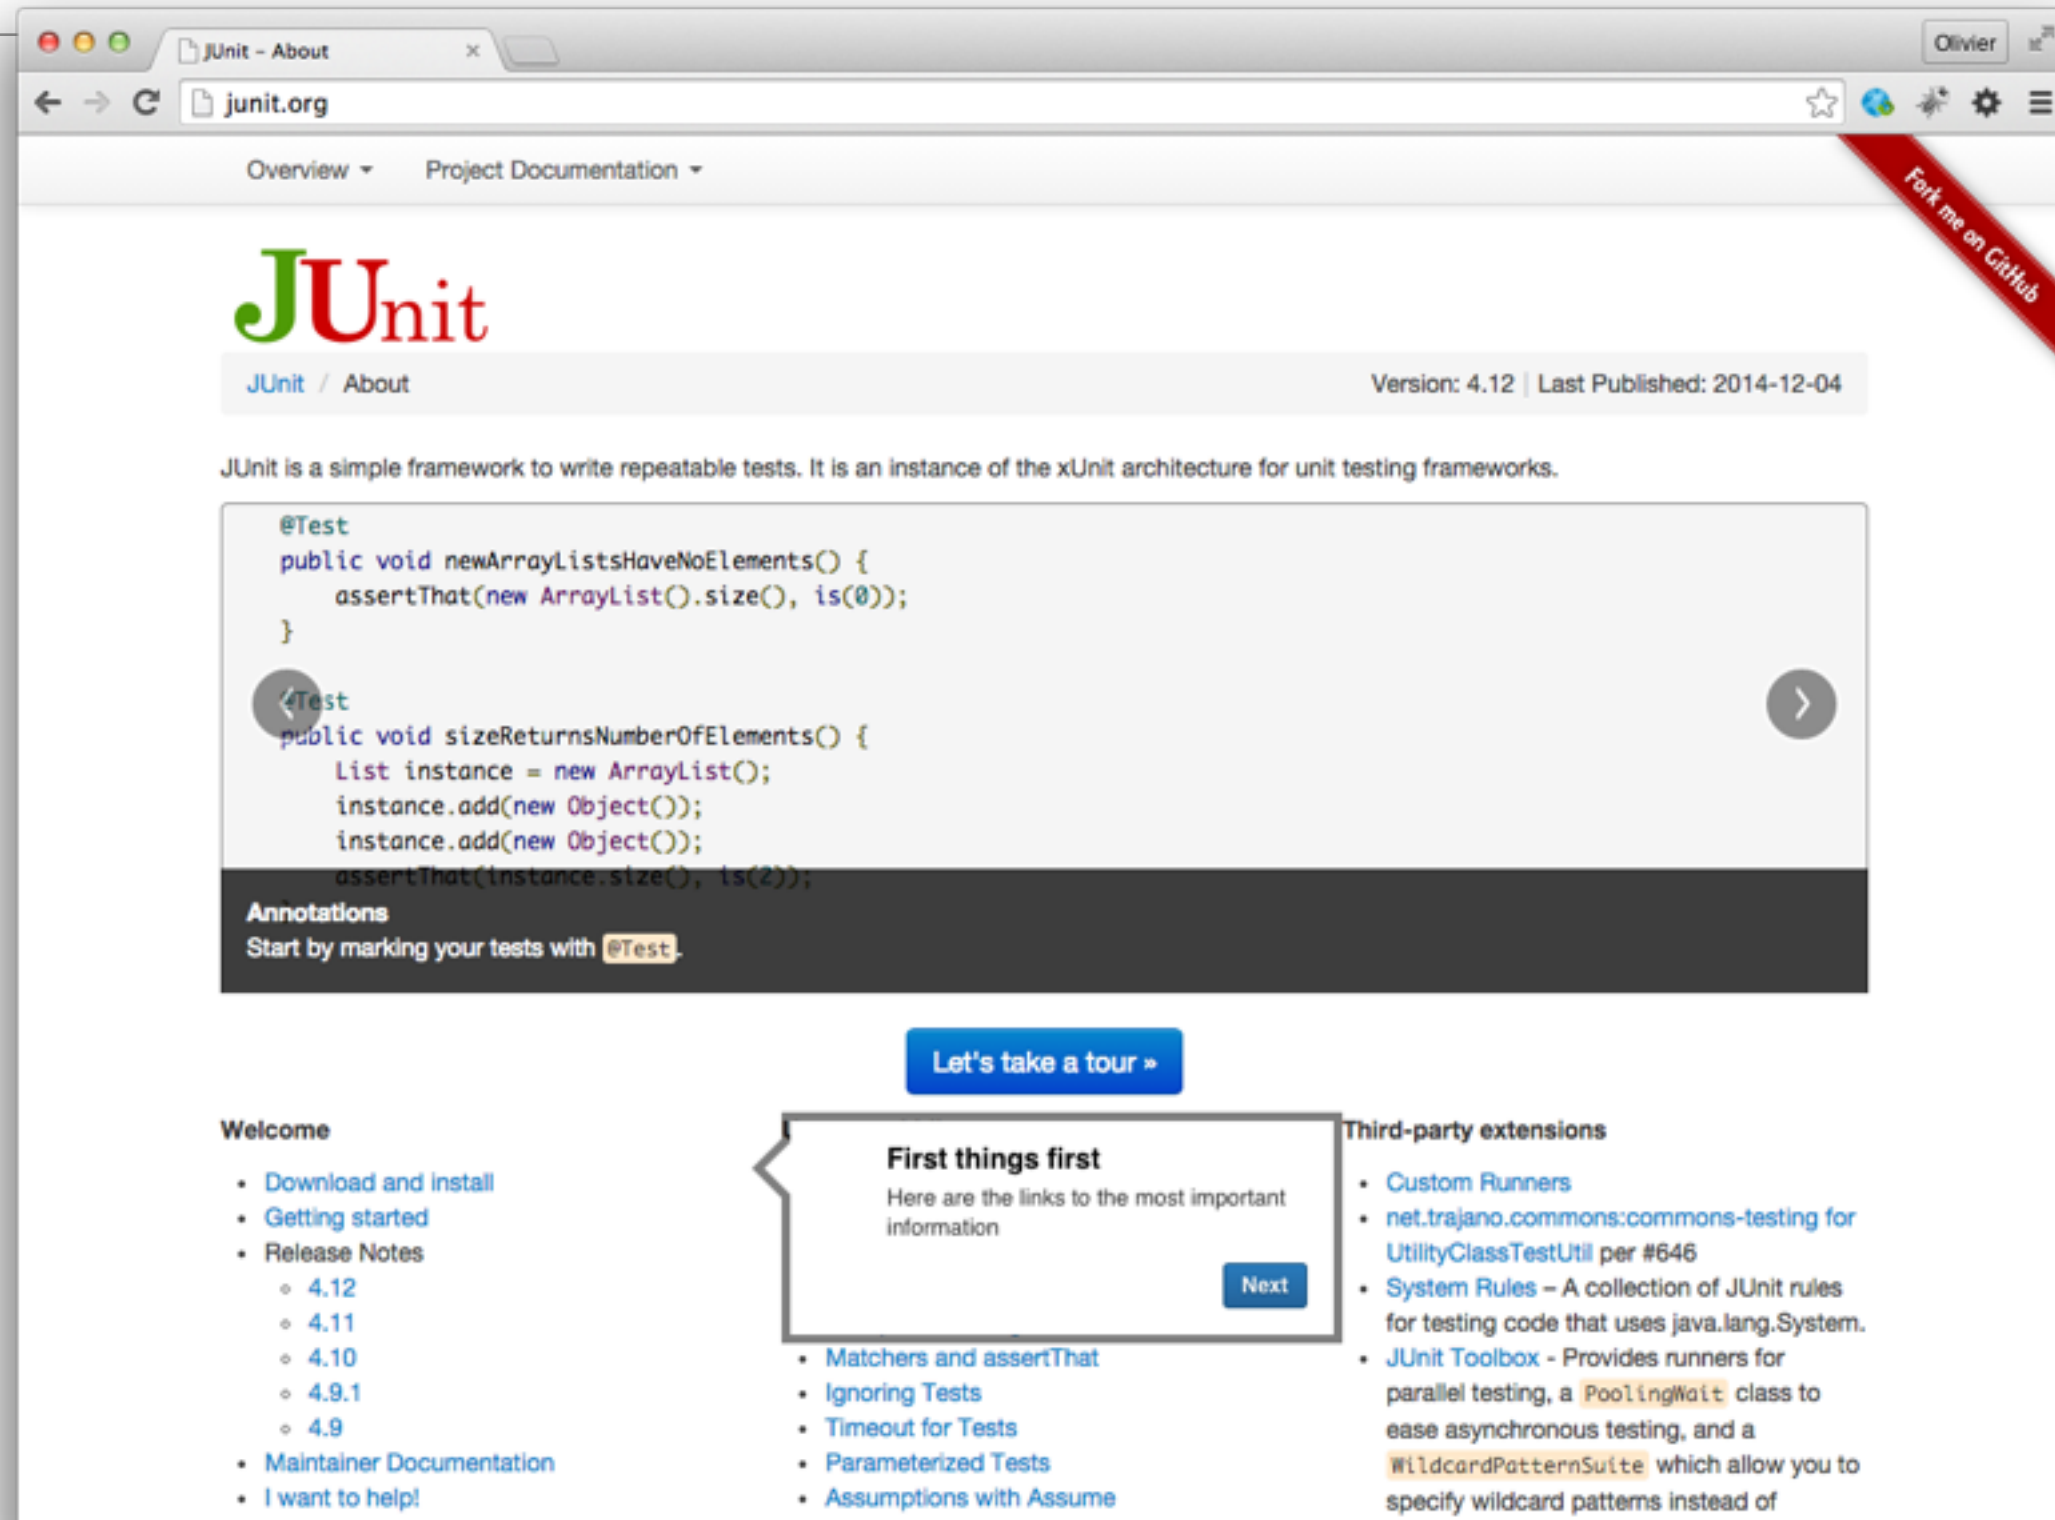
```
public interface IRouletteV2Client extends IRouletteV1Client {  
  
    public void clearDataStore() throws IOException;  
    public List<Student> listStudents() throws IOException;  
  
}
```



# Writing JUnit tests

heig-vd

Haute Ecole d'Ingénierie et de Gestion  
du Canton de Vaud



JUnit - About

junit.org

Overview ▾ Project Documentation ▾

# JUnit

JUnit / About

Version: 4.12 | Last Published: 2014-12-04

JUnit is a simple framework to write repeatable tests. It is an instance of the xUnit architecture for unit testing frameworks.

```
@Test
public void newArrayListsHaveNoElements() {
    assertThat(new ArrayList().size(), is(0));
}

@Test
public void sizeReturnsNumberOfElements() {
    List instance = new ArrayList();
    instance.add(new Object());
    instance.add(new Object());
    assertThat(instance.size(), is(2));
}
```

**Annotations**  
Start by marking your tests with `@Test`.

[Let's take a tour »](#)

## Welcome

- [Download and install](#)
- [Getting started](#)
- [Release Notes](#)
  - [4.12](#)
  - [4.11](#)
  - [4.10](#)
  - [4.9.1](#)
  - [4.9](#)
- [Maintainer Documentation](#)
- [I want to help!](#)

## First things first

Here are the links to the most important information

[Next](#)

- [Matchers and assertThat](#)
- [Ignoring Tests](#)
- [Timeout for Tests](#)
- [Parameterized Tests](#)
- [Assumptions with Assume](#)

## Third-party extensions

- [Custom Runners](#)
- [net.trajano.commons:commons-testing for UtilityClassTestUtil](#) per #646
- [System Rules](#) - A collection of JUnit rules for testing code that uses `java.lang.System`.
- [JUnit Toolbox](#) - Provides runners for parallel testing, a `PoolingWait` class to ease asynchronous testing, and a `WildcardPatternSuite` which allow you to specify wildcard patterns instead of

# Writing JUnit tests

**If we want to check that our code throws an exception  
in certain conditions, we can use this object**

```
public class RouletteV1WasadigiTest {  
  
    @Rule  
    public ExpectedException exception = ExpectedException.none();  
  
    @Rule  
    public EphemeralClientServerPair roulettePair =  
        new EphemeralClientServerPair(RouletteV1Protocol.VERSION);  
  
    @Test  
    @TestAuthor(githubId = "wasadigi")  
    public void theTestRouletteServerShouldRunDuringTests() throws IOException {  
        assertTrue(roulettePair.getServer().isRunning());  
    }  
    ...  
}
```

**For every test in the class, we  
want to start a Roulette Server,  
create a Roulette Client and  
ask the client to establish a  
connection with the server.**

***We don't want to duplicate this  
code again and again.***

**@Test indicates that this is a test method**

**JUnit provides methods to make assertions  
(check that certain conditions are met)**

**We have defined a  
custom annotation that  
we use to collect  
statistics about test  
authors**

# Writing JUnit tests

```
public class EphemeralClientServerPair extends ExternalResource {
```

```
...
```

```
@Override
```

```
protected void before() throws Throwable {
```

```
    server = new RouletteServer(protocolVersion);
```

```
    server.startServer();
```

```
    if (RouletteV1Protocol.VERSION.equals(protocolVersion)) {
```

```
        client = new RouletteV1ClientImpl();
```

```
    } else {
```

```
        client = new RouletteV2ClientImpl();
```

```
    }
```

```
    client.connect("localhost", server.getPort());
```

```
}
```

**This is executed before each test.**



**This is executed after each test.**



```
@Override
```

```
protected void after() {
```

```
    try {
```

```
        client.disconnect();
```

```
    } catch (IOException ex) {
```

```
...
```

```
    }
```

```
    try {
```

```
        server.stopServer();
```

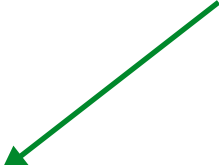
```
    } catch (IOException ex) {
```

```
...
```

```
    }
```

```
}
```

**Use these getters to  
access the client and the  
server from your test  
methods**



```
public RouletteServer getServer() {  
    return server;  
}
```

```
public IRouletteV1Client getClient() {  
    return client;  
}
```

# Writing JUnit tests


---

```
@Test
@TestAuthor(githubId = {"wasadigi", "SoftEng-HEIGVD"})
public void theServerShouldStillHaveZeroStudentsAtStart() throws IOException {
    assertEquals(0, roulettePair.getClient().getNumberOfStudents());
}
```

```
@Test
@TestAuthor(githubId = "SoftEng-HEIGVD")
public void theServerShouldCountStudents() throws IOException {
    IRouletteV1Client client = roulettePair.getClient();
    assertEquals(0, client.getNumberOfStudents());
    client.loadStudent("sacha");
    assertEquals(1, client.getNumberOfStudents());
    client.loadStudent("olivier");
    assertEquals(2, client.getNumberOfStudents());
    client.loadStudent("fabienne");
    assertEquals(3, client.getNumberOfStudents());
}
```

# Workflow: Step 1

- The first thing that we ask you to do is to **add new JUnit tests** to your project.
- You will do that in **a new branch** named **fb-lab02-tests** (git checkout -b)
- You will **submit a pull request** and **we will merge** your contribution on master.



```
$ git checkout -b fb-lab02-tests
$ git add ...
$ git commit -m "adding junit tests"
$ git push origin fb-lab02-tests
```

# Workflow: Step 2

---

- When we have received all contributions from all students, **you will pull them in your local clone**. Hence, you will have a project with a lot of JUnit class files (one per student or per group).
- Since all students create new files, you should have **no merge conflicts**.

**Every group should write at least 6 tests for the Protocol v1 and 6 tests for the Protocol v2.**

```
$ git checkout master  
$ git pull --rebase  
$ mvn clean install
```

# Workflow: Step 3

---

- You will **implement the application code**:
  - Implement the Roulette Protocol Client (v1 and v2)
  - Implement the Roulette Protocol Server (v2)